

MLOPS Group 114 - EDA analysis notebook

SNo	Name	ID	Percentage
1	Ravindra Babu Katiki	2024AB05286	100
2	Deepak Kumaran	2024AB05107	100
3	Degala Venkatesh	2023AC05236	100
4	Hemant Dyavarkonda	2024AB05109	100
5	Ashish Kumar	2024AB05110	100

TASK 1: Data Acquisition & Exploratory Data Analysis (EDA)

This notebook is exploratory and marked optional in assignment, but is provided as a supplement to the existing make suite that is the backbone of the overall experiment

#1.1 Load dependent libraries

```
import sys

!python -m pip install -U pip setuptools wheel
!python -m pip install "numpy>=2.1.0" "pandas>=2.2.3" "scikit-learn>=1.5.0" "matplotlib>=3.9.0" "joblib>=1.4.2"
!python -m pip install "mlflow>=2.12.2" "fastapi>=0.110.2"
"uvicorn>=0.29.0" "pydantic>=2.7.1"

Defaulting to user installation because normal site-packages is not
writeable
Requirement already satisfied: pip in c:\users\ravindra\appdata\roaming\python\python313\site-packages (25.3)
Requirement already satisfied: setuptools in c:\users\ravindra\appdata\roaming\python\python313\site-packages (80.9.0)
Requirement already satisfied: wheel in c:\users\ravindra\appdata\roaming\python\python313\site-packages (0.45.1)
Defaulting to user installation because normal site-packages is not
writeable
Requirement already satisfied: numpy>=2.1.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (2.3.5)
Requirement already satisfied: pandas>=2.2.3 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (2.3.3)
Requirement already satisfied: scikit-learn>=1.5.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (1.8.0)
Requirement already satisfied: matplotlib>=3.9.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (3.10.7)
Requirement already satisfied: joblib>=1.4.2 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (1.5.2)
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\
```

```
ravindra\appdata\roaming\python\python313\site-packages (from
pandas>=2.2.3) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from pandas>=2.2.3)
(2025.2)
Requirement already satisfied: tzdata>=2022.7 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from pandas>=2.2.3)
(2025.2)
Requirement already satisfied: scipy>=1.10.0 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from scikit-
learn>=1.5.0) (1.16.3)
Requirement already satisfied: threadpoolctl>=3.2.0 in c:\users\
ravindra\appdata\roaming\python\python313\site-packages (from scikit-
learn>=1.5.0) (3.6.0)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from
matplotlib>=3.9.0) (1.3.3)
Requirement already satisfied: cycler>=0.10 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from
matplotlib>=3.9.0) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from
matplotlib>=3.9.0) (4.61.0)
Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from
matplotlib>=3.9.0) (1.4.9)
Requirement already satisfied: packaging>=20.0 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from
matplotlib>=3.9.0) (25.0)
Requirement already satisfied: pillow>=8 in c:\users\ravindra\appdata\
roaming\python\python313\site-packages (from matplotlib>=3.9.0)
(12.0.0)
Requirement already satisfied: pyparsing>=3 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from
matplotlib>=3.9.0) (3.2.5)
Requirement already satisfied: six>=1.5 in c:\users\ravindra\appdata\
roaming\python\python313\site-packages (from python-dateutil>=2.8.2-
>pandas>=2.2.3) (1.17.0)
Defaulting to user installation because normal site-packages is not
writeable
Requirement already satisfied: mlflow>=2.12.2 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (3.8.0)
Requirement already satisfied: fastapi>=0.110.2 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (0.127.0)
Requirement already satisfied: uvicorn>=0.29.0 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (0.40.0)
Requirement already satisfied: pydantic>=2.7.1 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (2.12.5)
Requirement already satisfied: mlflow-skinny==3.8.0 in c:\users\
```

```
ravindra\appdata\roaming\python\python313\site-packages (from
mlflow>=2.12.2) (3.8.0)
Requirement already satisfied: mlflow-tracing==3.8.0 in c:\users\
ravindra\appdata\roaming\python\python313\site-packages (from
mlflow>=2.12.2) (3.8.0)
Requirement already satisfied: Flask-CORS<7 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from mlflow>=2.12.2)
(6.0.2)
Requirement already satisfied: Flask<4 in c:\users\ravindra\appdata\
roaming\python\python313\site-packages (from mlflow>=2.12.2) (3.1.2)
Requirement already satisfied: alembic!=1.10.0,<2 in c:\users\
ravindra\appdata\roaming\python\python313\site-packages (from
mlflow>=2.12.2) (1.17.2)
Requirement already satisfied: cryptography<47,>=43.0.0 in c:\users\
ravindra\appdata\roaming\python\python313\site-packages (from
mlflow>=2.12.2) (46.0.3)
Requirement already satisfied: docker<8,>=4.0.0 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from mlflow>=2.12.2)
(7.1.0)
Requirement already satisfied: graphene<4 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from mlflow>=2.12.2)
(3.4.3)
Requirement already satisfied: huey<3,>=2.5.0 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from mlflow>=2.12.2)
(2.5.5)
Requirement already satisfied: matplotlib<4 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from mlflow>=2.12.2)
(3.10.7)
Requirement already satisfied: numpy<3 in c:\users\ravindra\appdata\
roaming\python\python313\site-packages (from mlflow>=2.12.2) (2.3.5)
Requirement already satisfied: pandas<3 in c:\users\ravindra\appdata\
roaming\python\python313\site-packages (from mlflow>=2.12.2) (2.3.3)
Requirement already satisfied: pyarrow<23,>=4.0.0 in c:\users\
ravindra\appdata\roaming\python\python313\site-packages (from
mlflow>=2.12.2) (22.0.0)
Requirement already satisfied: scikit-learn<2 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from mlflow>=2.12.2)
(1.8.0)
Requirement already satisfied: scipy<2 in c:\users\ravindra\appdata\
roaming\python\python313\site-packages (from mlflow>=2.12.2) (1.16.3)
Requirement already satisfied: sqlalchemy<3,>=1.4.0 in c:\users\
ravindra\appdata\roaming\python\python313\site-packages (from
mlflow>=2.12.2) (2.0.45)
Requirement already satisfied: waitress<4 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from mlflow>=2.12.2)
(3.0.2)
Requirement already satisfied: cachetools<7,>=5.0.0 in c:\users\
ravindra\appdata\roaming\python\python313\site-packages (from mlflow-
skinny==3.8.0->mlflow>=2.12.2) (6.2.4)
```

```
Requirement already satisfied: click<9,>=7.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (8.3.1)
Requirement already satisfied: cloudpickle<4 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (3.1.2)
Requirement already satisfied: databricks-sdk<1,>=0.20.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (0.76.0)
Requirement already satisfied: gitpython<4,>=3.1.9 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (3.1.45)
Requirement already satisfied: importlib_metadata!=4.7.0,<9,>=3.7.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (8.7.1)
Requirement already satisfied: opentelemetry-api<3,>=1.9.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (1.39.1)
Requirement already satisfied: opentelemetry-proto<3,>=1.9.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (1.39.1)
Requirement already satisfied: opentelemetry-sdk<3,>=1.9.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (1.39.1)
Requirement already satisfied: packaging<26 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (25.0)
Requirement already satisfied: protobuf<7,>=3.12.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (6.33.2)
Requirement already satisfied: python-dotenv<2,>=0.19.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (1.2.1)
Requirement already satisfied: pyyaml<7,>=5.1 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (6.0.3)
Requirement already satisfied: requests<3,>=2.17.3 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (2.32.5)
Requirement already satisfied: sqlparse<1,>=0.4.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (0.5.5)
Requirement already satisfied: typing-extensions<5,>=4.0.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from mlflow-skinny==3.8.0->mlflow>=2.12.2) (4.15.0)
Requirement already satisfied: starlette<0.51.0,>=0.40.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from fastapi>=0.110.2) (0.50.0)
Requirement already satisfied: annotated-doc>=0.0.2 in c:\users\
```

```
ravindra\appdata\roaming\python\python313\site-packages (from
fastapi>=0.110.2) (0.0.4)
Requirement already satisfied: h11>=0.8 in c:\users\ravindra\appdata\
roaming\python\python313\site-packages (from uvicorn>=0.29.0) (0.16.0)
Requirement already satisfied: annotated-types>=0.6.0 in c:\users\
ravindra\appdata\roaming\python\python313\site-packages (from
pydantic>=2.7.1) (0.7.0)
Requirement already satisfied: pydantic-core==2.41.5 in c:\users\
ravindra\appdata\roaming\python\python313\site-packages (from
pydantic>=2.7.1) (2.41.5)
Requirement already satisfied: typing-inspection>=0.4.2 in c:\users\
ravindra\appdata\roaming\python\python313\site-packages (from
pydantic>=2.7.1) (0.4.2)
Requirement already satisfied: Mako in c:\users\ravindra\appdata\
roaming\python\python313\site-packages (from alembic!=1.10.0,<2-
>mlflow>=2.12.2) (1.3.10)
Requirement already satisfied: colorama in c:\users\ravindra\appdata\
roaming\python\python313\site-packages (from click<9,>=7.0->mlflow-
skinny==3.8.0->mlflow>=2.12.2) (0.4.6)
Requirement already satisfied: cffi>=2.0.0 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from
cryptography<47,>=43.0.0->mlflow>=2.12.2) (2.0.0)
Requirement already satisfied: google-auth~=2.0 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from databricks-
sdk<1,>=0.20.0->mlflow-skinny==3.8.0->mlflow>=2.12.2) (2.45.0)
Requirement already satisfied: pywin32>=304 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from docker<8,>=4.0.0-
>mlflow>=2.12.2) (311)
Requirement already satisfied: urllib3>=1.26.0 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from docker<8,>=4.0.0-
>mlflow>=2.12.2) (2.5.0)
Requirement already satisfied: blinker>=1.9.0 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from Flask<4-
>mlflow>=2.12.2) (1.9.0)
Requirement already satisfied: itsdangerous>=2.2.0 in c:\users\
ravindra\appdata\roaming\python\python313\site-packages (from Flask<4-
>mlflow>=2.12.2) (2.2.0)
Requirement already satisfied: jinja2>=3.1.2 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from Flask<4-
>mlflow>=2.12.2) (3.1.6)
Requirement already satisfied: markupsafe>=2.1.1 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from Flask<4-
>mlflow>=2.12.2) (3.0.3)
Requirement already satisfied: werkzeug>=3.1.0 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from Flask<4-
>mlflow>=2.12.2) (3.1.4)
Requirement already satisfied: gitdb<5,>=4.0.1 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from
gitpython<4,>=3.1.9->mlflow-skinny==3.8.0->mlflow>=2.12.2) (4.0.12)
```

```
Requirement already satisfied: smmap<6,>=3.0.1 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from gitdb<5,>=4.0.1->gitpython<4,>=3.1.9->mlflow-skinny==3.8.0->mlflow>=2.12.2) (5.0.2)
Requirement already satisfied: pyasn1-modules>=0.2.1 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from google-auth~=2.0->databricks-sdk<1,>=0.20.0->mlflow-skinny==3.8.0->mlflow>=2.12.2) (0.4.2)
Requirement already satisfied: rsa<5,>=3.1.4 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from google-auth~=2.0->databricks-sdk<1,>=0.20.0->mlflow-skinny==3.8.0->mlflow>=2.12.2) (4.9.1)
Requirement already satisfied: graphql-core<3.3,>=3.1 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from graphene<4->mlflow>=2.12.2) (3.2.7)
Requirement already satisfied: graphql-relay<3.3,>=3.1 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from graphene<4->mlflow>=2.12.2) (3.2.0)
Requirement already satisfied: python-dateutil<3,>=2.7.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from graphene<4->mlflow>=2.12.2) (2.9.0.post0)
Requirement already satisfied: zipp>=3.20 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from importlib_metadata!=4.7.0,<9,>=3.7.0->mlflow-skinny==3.8.0->mlflow>=2.12.2) (3.23.0)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from matplotlib<4->mlflow>=2.12.2) (1.3.3)
Requirement already satisfied: cycler>=0.10 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from matplotlib<4->mlflow>=2.12.2) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from matplotlib<4->mlflow>=2.12.2) (4.61.0)
Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from matplotlib<4->mlflow>=2.12.2) (1.4.9)
Requirement already satisfied: pillow>=8 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from matplotlib<4->mlflow>=2.12.2) (12.0.0)
Requirement already satisfied: pyparsing>=3 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from matplotlib<4->mlflow>=2.12.2) (3.2.5)
Requirement already satisfied: opentelemetry-semantic-conventions==0.60b1 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from opentelemetry-sdk<3,>=1.9.0->mlflow-skinny==3.8.0->mlflow>=2.12.2) (0.60b1)
Requirement already satisfied: pytz>=2020.1 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from pandas<3->mlflow>=2.12.2) (2025.2)
```

```
Requirement already satisfied: tzdata>=2022.7 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from pandas<3->mlflow>=2.12.2) (2025.2)
Requirement already satisfied: six>=1.5 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from python-dateutil<3,>=2.7.0->graphene<4->mlflow>=2.12.2) (1.17.0)
Requirement already satisfied: charset_normalizer<4,>=2 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from requests<3,>=2.17.3->mlflow-skinny==3.8.0->mlflow>=2.12.2) (3.4.4)
Requirement already satisfied: idna<4,>=2.5 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from requests<3,>=2.17.3->mlflow-skinny==3.8.0->mlflow>=2.12.2) (3.11)
Requirement already satisfied: certifi>=2017.4.17 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from requests<3,>=2.17.3->mlflow-skinny==3.8.0->mlflow>=2.12.2) (2025.11.12)
Requirement already satisfied: pyasn1>=0.1.3 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from rsa<5,>=3.1.4->google-auth~2.0->databricks-sdk<1,>=0.20.0->mlflow-skinny==3.8.0->mlflow>=2.12.2) (0.6.1)
Requirement already satisfied: joblib>=1.3.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from scikit-learn<2->mlflow>=2.12.2) (1.5.2)
Requirement already satisfied: threadpoolctl>=3.2.0 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from scikit-learn<2->mlflow>=2.12.2) (3.6.0)
Requirement already satisfied: greenlet>=1 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from sqlalchemy<3,>=1.4.0->mlflow>=2.12.2) (3.3.0)
Requirement already satisfied: anyio<5,>=3.6.2 in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from starlette<0.51.0,>=0.40.0->fastapi>=0.110.2) (4.12.0)
Requirement already satisfied: pycparser in c:\users\ravindra\appdata\roaming\python\python313\site-packages (from cffi>=2.0.0->cryptography<47,>=43.0.0->mlflow>=2.12.2) (2.23)
```

1.2 UCI Heart Dataset (Cardio Risk Ops)

```
import zipfile, urllib.request, io
import pandas as pd
from pathlib import Path

# Paths
PROJECT_ROOT = Path(".").resolve()
DATA_DIR = PROJECT_ROOT / "data" / "raw"
DATA_DIR.mkdir(parents=True, exist_ok=True)
CSV_PATH = DATA_DIR / "heart.csv"

# Official UCI ZIP
UCI_ZIP_URL =
```

```

"https://archive.ics.uci.edu/static/public/45/heart+disease.zip"

# Download ZIP
resp = urllib.request.urlopen(UCI_ZIP_URL)
zip_bytes = io.BytesIO(resp.read())

# Open ZIP and locate Cleveland file safely
with zipfile.ZipFile(zip_bytes) as z:
    names = z.namelist()

    cleveland_files = [n for n in names if "processed.cleveland.data"
in n]
    if not cleveland_files:
        raise RuntimeError("processed.cleveland.data not found in
ZIP")

    cleveland_path = cleveland_files[0]
    print("Using file from ZIP:", cleveland_path)

    with z.open(cleveland_path) as f:
        df = pd.read_csv(f, header=None)

# Assign official column names
df.columns = [
    "age", "sex", "cp", "trestbps", "chol", "fbs",
    "restecg", "thalach", "exang", "oldpeak",
    "slope", "ca", "thal", "target"
]

# Clean missing values
df.replace("?", pd.NA, inplace=True)
df = df.apply(pd.to_numeric, errors="coerce")

# Convert target to binary
df["target"] = df["target"].apply(lambda x: 1 if x > 0 else 0)

# Save CSV
df.to_csv(CSV_PATH, index=False)

print("Saved CSV to:", CSV_PATH)
df.head()

Using file from ZIP: processed.cleveland.data
Saved CSV to: C:\Users\Ravindra\Documents\cardio_risk_ops\data\raw\
heart.csv

      age  sex   cp  trestbps   chol   fbs  restecg  thalach  exang
oldpeak \
0    63.0  1.0  1.0       145.0  233.0  1.0        2.0     150.0     0.0
2.3

```

1	67.0	1.0	4.0	160.0	286.0	0.0	2.0	108.0	1.0
1.5									
2	67.0	1.0	4.0	120.0	229.0	0.0	2.0	129.0	1.0
2.6									
3	37.0	1.0	3.0	130.0	250.0	0.0	0.0	187.0	0.0
3.5									
4	41.0	0.0	2.0	130.0	204.0	0.0	2.0	172.0	0.0
1.4									
	slope	ca	thal	target					
0	3.0	0.0	6.0	0					
1	2.0	3.0	3.0	1					
2	2.0	2.0	7.0	1					
3	3.0	0.0	3.0	0					
4	1.0	0.0	3.0	0					

1) Basic shape and schema

We first check the number of rows/columns and the feature names to confirm the dataset is loaded correctly.

```
df.shape, list(df.columns)

((303, 14),
 ['age',
 'sex',
 'cp',
 'trestbps',
 'chol',
 'fbs',
 'restecg',
 'thalach',
 'exang',
 'oldpeak',
 'slope',
 'ca',
 'thal',
 'target'])
```

2) Missing values

A quick missing-value summary helps decide whether we need imputation.


```
missing = df.isna().sum().sort_values(ascending=False)
missing[missing > 0].head(20)
```

```
ca      4  
thal    2  
dtype: int64
```

3) Target distribution (class balance)

If the dataset is imbalanced, accuracy can be misleading. We will also report ROC-AUC during modeling.

```
import matplotlib.pyplot as plt  
  
ax = df["target"].value_counts().sort_index().plot(kind="bar")  
ax.set_title("Class Balance (target)")  
ax.set_xlabel("Class (0 = no disease, 1 = disease)")  
ax.set_ylabel("Count")  
plt.tight_layout()  
plt.show()
```



4) Summary statistics (numeric)

This shows typical ranges and potential outliers for key continuous variables.

```
df.describe().T
```


	count	mean	std	min	25%	50%	75%
max							
age	303.0	54.438944	9.038662	29.0	48.0	56.0	61.0
77.0							
sex	303.0	0.679868	0.467299	0.0	0.0	1.0	1.0
1.0							
cp	303.0	3.158416	0.960126	1.0	3.0	3.0	4.0
4.0							
trestbps	303.0	131.689769	17.599748	94.0	120.0	130.0	140.0
200.0							
chol	303.0	246.693069	51.776918	126.0	211.0	241.0	275.0
564.0							
fbs	303.0	0.148515	0.356198	0.0	0.0	0.0	0.0
1.0							
restecg	303.0	0.990099	0.994971	0.0	0.0	1.0	2.0
2.0							
thalach	303.0	149.607261	22.875003	71.0	133.5	153.0	166.0
202.0							
exang	303.0	0.326733	0.469794	0.0	0.0	0.0	1.0
1.0							
oldpeak	303.0	1.039604	1.161075	0.0	0.0	0.8	1.6
6.2							
slope	303.0	1.600660	0.616226	1.0	1.0	2.0	2.0
3.0							
ca	299.0	0.672241	0.937438	0.0	0.0	0.0	1.0
3.0							
thal	301.0	4.734219	1.939706	3.0	3.0	3.0	7.0
7.0							
target	303.0	0.458746	0.499120	0.0	0.0	0.0	1.0
1.0							

5) Distributions of numeric features

We plot histograms for key numeric features to understand their spread.

```
num_cols = ["age", "trestbps", "chol", "thalach", "oldpeak"]
df[num_cols].hist(bins=20, figsize=(10, 6))
plt.suptitle("Numeric Feature Distributions", y=1.02)
plt.tight_layout()
plt.show()
```

Numeric Feature Distributions




6) Correlation heatmap (numeric only)

Correlation can hint at relationships among data attributes (features)

```
corr = df.corr(numeric_only=True)

plt.figure(figsize=(8, 6))
plt.imshow(corr, aspect="auto")
plt.colorbar()
plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
plt.yticks(range(len(corr.columns)), corr.columns)
plt.title("Correlation Heatmap (numeric)")
plt.tight_layout()
plt.show()
```



7) Simple feature-vs-target view using boxplots

Boxplots are a quick way to compare distributions between classes.

```


fig, axes = plt.subplots(2, 3, figsize=(12, 7))
axes = axes.flatten()

for i, col in enumerate(num_cols):
    df.boxplot(column=col, by="target", ax=axes[i])
    axes[i].set_title(col)
    axes[i].set_xlabel("target")
    axes[i].set_ylabel(col)

# Remove extra subplot if any
for j in range(len(num_cols), len(axes)):
    fig.delaxes(axes[j])

plt.suptitle("")
plt.tight_layout()
plt.show()

```



Task 2. Feature Engineering & Model Development

Modeling, Evaluation, and Inference (Extension)

This section extends the notebook from **EDA** to a complete **end-to-end** workflow:

- Train/test split (stratified)
- Reproducible preprocessing + model pipelines
- Two models (Logistic Regression baseline + tuned Random Forest)
- Metrics + ROC curve + confusion matrix
- Lightweight explainability (permutation importance)
- Save model artifact and reload
- Inference test on a single payload (production-style)

```
# Modeling imports
import sys

!{sys.executable} -m pip install scikit-learn

from sklearn.model_selection import train_test_split
from sklearn.pipeline import Pipeline
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, roc_auc_score,
    classification_report, confusion_matrix
)
from sklearn.linear_model import LogisticRegression
```

```

from sklearn.ensemble import RandomForestClassifier

import numpy as np

Defaulting to user installation because normal site-packages is not
writeable
Requirement already satisfied: scikit-learn in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (1.8.0)
Requirement already satisfied: numpy>=1.24.1 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from scikit-learn)
(2.3.5)
Requirement already satisfied: scipy>=1.10.0 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from scikit-learn)
(1.16.3)
Requirement already satisfied: joblib>=1.3.0 in c:\users\ravindra\
appdata\roaming\python\python313\site-packages (from scikit-learn)
(1.5.2)
Requirement already satisfied: threadpoolctl>=3.2.0 in c:\users\
ravindra\appdata\roaming\python\python313\site-packages (from scikit-
learn) (3.6.0)

# print src path to load cardiooops

import sys
from pathlib import Path

PROJECT_ROOT = Path("..").resolve()
SRC_PATH = PROJECT_ROOT / "src"

if str(SRC_PATH) not in sys.path:
    sys.path.insert(0, str(SRC_PATH))

print("SRC_PATH added:", SRC_PATH)

SRC_PATH added: C:\Users\Ravindra\Documents\cardio_risk_ops\src

assert 'df' in globals(), "df is not loaded. Please run the data
loading cells above first."
df.shape, df.columns.tolist()[:5]

((303, 14), ['age', 'sex', 'cp', 'trestbps', 'chol'])

```

1) Train / test split (stratified)

```

# Split data

from sklearn.model_selection import train_test_split

X = df.drop(columns=["target"])
y = df["target"].astype(int)

```

```

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    stratify=y,
    random_state=42
)

X_train.shape, X_test.shape, y_train.mean(), y_test.mean()

((242, 13),
(61, 13),
np.float64(0.45867768595041325),
np.float64(0.45901639344262296))

```

2) Preprocessing pipeline (same for training and inference)

```

from cardioops.features import build_preprocessor

preprocessor = build_preprocessor()
preprocessor

ColumnTransformer(sparse_threshold=0.2,
                  transformers=[('num',
                                 Pipeline(steps=[('impute',
                                                 SimpleImputer(strategy='median')),
                                                ('scale',
                                                 StandardScaler()))),
                               ('cat',
                                 Pipeline(steps=[('impute',
                                                 SimpleImputer(strategy='most_frequent')),
                                                ('onehot',
                                                 OneHotEncoder(handle_unknown='ignore'))])),
                               ('oldpeak',
                                 Pipeline(steps=[('impute',
                                                 SimpleImputer(strategy='median')),
                                                ('scale',
                                                 StandardScaler())]))],
                  sparse_threshold=0.2)

```

3) Model 1 — Logistic Regression (baseline)

Baseline models are useful because they are fast, interpretable, and provide a reliable reference ROC-AUC.

```

from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression

logreg = Pipeline(steps=[
    ("prep", preprocessor),
    ("clf", LogisticRegression(max_iter=2500, solver="liblinear",
random_state=42))
])

logreg.fit(X_train, y_train)

Pipeline(steps=[('prep',
                 ColumnTransformer(sparse_threshold=0.2,
                                     transformers=[('num',
                                     Pipeline(steps=[('impute',
                                                     SimpleImputer(strategy='median'))),
                                     ('scale',
                                      StandardScaler()))),
                                     ['age', 'trestbps',
                                      'chol',
                                      'oldpeak']),
                 ('cat',
                  Pipeline(steps=[('impute',
                                  SimpleImputer(strategy='most_frequent')),
                                 ('onehot',
                                  OneHotEncoder(handle_unknown='ignore'))]),
                  ['sex', 'cp', 'fbs',
                   'restecg', 'exang',
                   'slope',
                   ('clf',
                    LogisticRegression(max_iter=2500, random_state=42,
                                       solver='liblinear'))])),
                 # metrics for evaluation of Logistic regression
import matplotlib.pyplot as plt
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, roc_auc_score,
    ConfusionMatrixDisplay, RocCurveDisplay, classification_report
)

```

```
proba_lr = logreg.predict_proba(X_test)[:, 1]
pred_lr = (proba_lr >= 0.5).astype(int)

lr_metrics = {
    "accuracy": accuracy_score(y_test, pred_lr),
    "precision": precision_score(y_test, pred_lr, zero_division=0),
    "recall": recall_score(y_test, pred_lr, zero_division=0),
    "roc_auc": roc_auc_score(y_test, proba_lr),
}
lr_metrics


{'accuracy': 0.8852459016393442,
 'precision': 0.8387096774193549,
 'recall': 0.9285714285714286,
 'roc_auc': 0.9664502164502166}


# Logistic Regression ROC curve and Confusion table

RocCurveDisplay.from_predictions(y_test, proba_lr)
plt.title("ROC Curve – Logistic Regression")
plt.show()

ConfusionMatrixDisplay.from_predictions(y_test, pred_lr)
plt.title("Confusion Matrix – Logistic Regression")
plt.show()

print(classification_report(y_test, pred_lr, digits=3))
```





	precision	recall	f1-score	support
0	0.933	0.848	0.889	33
1	0.839	0.929	0.881	28
accuracy			0.885	61
macro avg	0.886	0.889	0.885	61
weighted avg	0.890	0.885	0.885	61

4) Model 2 — Random Forest

Random Forest captures non-linear relationships. We do a ** grid search** for ROC-AUC.

```
# grid search
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import StratifiedKFold, GridSearchCV

rf_pipe = Pipeline(steps=[
    ("prep", preprocessor),
    ("clf", RandomForestClassifier(random_state=42))])
```

```

])
param_grid = {
    "clf_n_estimators": [200, 400],
    "clf_max_depth": [None, 5, 10],
    "clf_min_samples_leaf": [1, 2, 4],
}
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

rf_search = GridSearchCV(
    estimator=rf_pipe,
    param_grid=param_grid,
    scoring="roc_auc",
    cv=cv,
    n_jobs=-1,
    refit=True
)
rf_search.fit(X_train, y_train)

rf_search.best_params_, rf_search.best_score_
({'clf_max_depth': 5, 'clf_min_samples_leaf': 2,
'clf_n_estimators': 400},
 np.float64(0.8992601601297252))

# evaluate metrics for RandomForest algorithm

best_rf = rf_search.best_estimator_
proba_rf = best_rf.predict_proba(X_test)[:, 1]
pred_rf = (proba_rf >= 0.5).astype(int)

rf_metrics = {
    "accuracy": accuracy_score(y_test, pred_rf),
    "precision": precision_score(y_test, pred_rf, zero_division=0),
    "recall": recall_score(y_test, pred_rf, zero_division=0),
    "roc_auc": roc_auc_score(y_test, proba_rf),
}
rf_metrics
{'accuracy': 0.8524590163934426,
 'precision': 0.8064516129032258,
 'recall': 0.8928571428571429,
 'roc_auc': 0.9545454545454546}

# Random Forest ROC curve and Confusion table


RocCurveDisplay.from_predictions(y_test, proba_rf)
plt.title("ROC Curve – Random Forest ")


```

```
plt.show()

ConfusionMatrixDisplay.from_predictions(y_test, pred_rf)
plt.title("Confusion Matrix – Random Forest ")
plt.show()

print(classification_report(y_test, pred_rf, digits=3))
```





	precision	recall	f1-score	support
0	0.900	0.818	0.857	33
1	0.806	0.893	0.847	28
accuracy			0.852	61
macro avg	0.853	0.856	0.852	61
weighted avg	0.857	0.852	0.853	61

5) Model comparison table

```
# compare models
import pandas as pd

pd.DataFrame([lr_metrics, rf_metrics], index=["LogReg",
"RandForest"]).sort_values("roc_auc", ascending=False)

      accuracy  precision  recall  roc_auc
LogReg    0.885246   0.838710  0.928571  0.966450
RandForest 0.852459   0.806452  0.892857  0.954545
```

6) Lightweight explainability — permutation importance

This provides a simple explanation of which inputs most affect ROC-AUC.

```
# Correct feature names for permutation importance on raw inputs

import pandas as pd
from sklearn.inspection import permutation_importance

perm = permutation_importance(
    best_rf, X_test, y_test,
    n_repeats=15,
    random_state=42,
    scoring="roc_auc"
)

# IMPORTANT: permutation_importance is on original X_test columns (13)
feature_names = list(X_test.columns)


imp = pd.Series(perm.importances_mean, index=feature_names) \
    .sort_values(ascending=False) \
    .head(15)

imp

ca          0.037734
cp          0.037302
exang       0.019264
oldpeak     0.018254
thal         0.015296
thalach      0.007431
slope        0.005195
age          0.003824
restecg      0.001659
sex          0.000361
chol         0.000072
trestbps     -0.001010
fbs          -0.001804
dtype: float64

# ROC-AUC impact

ax = imp.iloc[::-1].plot(kind="barh", figsize=(8, 6))
ax.set_title("Top Permutation Importances (ROC-AUC impact)")
ax.set_xlabel("Mean importance decrease")
plt.tight_layout()
plt.show()
```



7) Save best model pipeline (artifact)

```
# save model to MODEL_PATH

import joblib
from pathlib import Path

ARTIFACT_DIR = Path("../artifacts")
ARTIFACT_DIR.mkdir(exist_ok=True)

MODEL_PATH = ARTIFACT_DIR / "final_pipeline.joblib"
joblib.dump(best_rf, MODEL_PATH)

MODEL_PATH.exists(), MODEL_PATH
(True, WindowsPath('../artifacts/final_pipeline.joblib'))
```

8) Reload model + inference test (production-style payload)

```
# Reload model

loaded_model = joblib.load(MODEL_PATH)

sample_payload = {
    "age": 63, "sex": 1, "cp": 3, "trestbps": 145, "chol": 233, "fbs": 1,
    "restecg": 0,
    "thalach": 150, "exang": 0, "oldpeak": 2.3, "slope": 0, "ca": 0,
    "thal": 1
}

sample_df = pd.DataFrame([sample_payload])

prob = float(loaded_model.predict_proba(sample_df)[:, 1][0])
label = int(prob >= 0.5)

{"risk_label": label, "risk_probability": round(prob, 6)}

{'risk_label': 0, 'risk_probability': 0.393604}
```

9) Threshold tuning

We evaluate a few thresholds to see precision/recall trade-offs.

```
# Threshold tuning

thresholds = [0.3, 0.4, 0.5, 0.6, 0.7]
rows = []
for t in thresholds:
    pred_t = (proba_rf >= t).astype(int)
    rows.append({
        "threshold": t,
        "precision": precision_score(y_test, pred_t, zero_division=0),
        "recall": recall_score(y_test, pred_t, zero_division=0),
        "accuracy": accuracy_score(y_test, pred_t),
    })
pd.DataFrame(rows)

   threshold  precision      recall  accuracy
0         0.3  0.736842  1.000000  0.836066
1         0.4  0.794118  0.964286  0.868852
2         0.5  0.806452  0.892857  0.852459
3         0.6  0.880000  0.785714  0.852459
4         0.7  0.947368  0.642857  0.819672
```

```
import sys
print(sys.version)

3.13.7 (tags/v3.13.7:bceelc3, Aug 14 2025, 14:15:11) [MSC v.1944 64
bit (AMD64)]
```

Task 3: Experiment Tracking using MLFlow

```
# Import MLFlow

import mlflow
import mlflow.sklearn
from mlflow.models.signature import infer_signature
from pathlib import Path

# -----
# Define numerical & categorical columns
# -----

# Target column
TARGET_COL = "target" # change only if your label column has a
different name

# All feature columns
feature_cols = [c for c in df.columns if c != TARGET_COL]

# Numeric columns (int / float)
num_cols = df[feature_cols].select_dtypes(include=["int64",
"float64"]).columns.tolist()

# Categorical columns (object / category / bool)
cat_cols = df[feature_cols].select_dtypes(
    include=["object", "category", "bool"])
).columns.tolist()

print("Numerical columns:", num_cols)
print("Categorical columns:", cat_cols)

Numerical columns: ['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs',
'restecg', 'thalach', 'exang', 'oldpeak', 'slope', 'ca', 'thal']
Categorical columns: []

from pathlib import Path
import mlflow

# Find project root (portable)
PROJECT_ROOT = Path.cwd().resolve()
for p in [PROJECT_ROOT] + list(PROJECT_ROOT.parents):
    if (p / "mlruns").exists() or (p / ".git").exists():
        PROJECT_ROOT = p
        break
```

```
MLRUNS_DIR = PROJECT_ROOT / "mlruns"
MLRUNS_DIR.mkdir(parents=True, exist_ok=True)
(MLRUNS_DIR / ".trash").mkdir(parents=True, exist_ok=True)

mlflow.set_tracking_uri(MLRUNS_DIR.as_uri())

EXPERIMENT_NAME = "ravindra_cardio_risk_mllops_assignment1"
mlflow.set_experiment(EXPERIMENT_NAME)

# If a previous run is still open in *this kernel*, close it
try:
    mlflow.end_run(status="KILLED")
except Exception:
    pass

print("Tracking URI:", mlflow.get_tracking_uri())
print("Experiment:", EXPERIMENT_NAME)
print("Active run:", mlflow.active_run())

Tracking URI:
file:///C:/Users/Ravindra/Documents/cardio_risk_ops/notebooks/mlruns
Experiment: ravindra_cardio_risk_mllops_assignment1
Active run: None

C:\Users\Ravindra\AppData\Roaming\Python\Python313\site-packages\
mlflow\tracking\_tracking_service\utils.py:177: FutureWarning: The
filesystem tracking backend (e.g., './mlruns') will be deprecated in
February 2026. Consider transitioning to a database backend (e.g.,
'sqlite:///mlflow.db') to take advantage of the latest MLflow
features. See https://github.com/mlflow/mlflow/issues/18534 for more
details and migration guidance. For migrating existing data,
https://github.com/mlflow/mlflow-export-import can be used.
    return FileStore(store_uri, store_uri)

# =====
# MLflow Experiment Tracking LogisticRegression
# =====

from pathlib import Path
import numpy as np
import mlflow
import mlflow.sklearn

from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import (
    roc_auc_score, accuracy_score, precision_score, recall_score,
```

```

f1_score, confusion_matrix
)

# -----
# 0) MLflow local tracking
# -----
PROJECT_ROOT = Path.cwd().resolve()
MLRUNS_DIR = PROJECT_ROOT / "mlruns"
(MLRUNS_DIR / ".trash").mkdir(parents=True, exist_ok=True)

mlflow.set_tracking_uri(MLRUNS_DIR.as_uri())
EXPERIMENT_NAME = "ravindra_cardio_risk_mlops_assignment1"
mlflow.set_experiment(EXPERIMENT_NAME)

# -----
# 1) Stop any active run (re-run safe)
# -----
if mlflow.active_run() is not None:
    mlflow.end_run(status="KILLED")

# -----
# 2) Build preprocessing + model pipeline
# -----
num_pipe = Pipeline(steps=[
    ("impute", SimpleImputer(strategy="median")),
    ("scale", StandardScaler())
])

cat_pipe = Pipeline(steps=[
    ("impute", SimpleImputer(strategy="most_frequent")),
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

preprocessor = ColumnTransformer(
    transformers=[
        ("num", num_pipe, num_cols),
        ("cat", cat_pipe, cat_cols),
    ],
    remainder="drop"
)

logreg_pipeline = Pipeline(steps=[
    ("prep", preprocessor),
    ("clf", LogisticRegression(max_iter=2000, solver="liblinear",
random_state=42))
])

# -----
# 3) Train + log to MLflow
# -----

```

```

run = mlflow.start_run(run_name="baseline_logreg_notebook")

try:
    mlflow.set_tag("model_family", "logistic_regression")
    mlflow.set_tag("dataset", "UCI Heart Disease (Cleveland)")
    mlflow.log_param("max_iter", 2000)
    mlflow.log_param("solver", "liblinear")

    logreg_pipeline.fit(X_train, y_train)

    y_prob = logreg_pipeline.predict_proba(X_test)[:, 1]
    y_pred = (y_prob >= 0.5).astype(int)

    auc = roc_auc_score(y_test, y_prob)
    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred, zero_division=0)
    rec = recall_score(y_test, y_pred, zero_division=0)
    f1 = f1_score(y_test, y_pred, zero_division=0)
    cm = confusion_matrix(y_test, y_pred)

    mlflow.log_metric("roc_auc", float(auc))
    mlflow.log_metric("accuracy", float(acc))
    mlflow.log_metric("precision", float(prec))
    mlflow.log_metric("recall", float(rec))
    mlflow.log_metric("f1", float(f1))

    # log confusion matrix as params (simple + visible)
    mlflow.log_param("cm_00", int(cm[0, 0]))
    mlflow.log_param("cm_01", int(cm[0, 1]))
    mlflow.log_param("cm_10", int(cm[1, 0]))
    mlflow.log_param("cm_11", int(cm[1, 1]))

    mlflow.sklearn.log_model(
        sk_model=logreg_pipeline,
        name="model",
        input_example=X_test.iloc[:3]
    )

    print("LogReg finished.")
    print("ROC-AUC:", auc)

except Exception:
    mlflow.end_run(status="FAILED")
    raise

finally:
    if mlflow.active_run() is not None:
        mlflow.end_run(status="FINISHED")

```

```
C:\Users\Ravindra\AppData\Roaming\Python\Python313\site-packages\
mlflow\types\utils.py:452: UserWarning: Hint: Inferred schema contains
integer column(s). Integer columns in Python cannot represent missing
values. If your input data contains missing values at inference time,
it will be encoded as floats and will cause a schema enforcement
error. The best way to avoid this problem is to infer the model schema
based on a realistic data sample (training dataset) that includes
missing values. Alternatively, you can declare integer columns as
doubles (float64) whenever these columns may have missing values. See
`Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-missing-values>`_ for more details.
    warnings.warn(
Downloading artifacts: 100%|██████████| 7/7
[00:00<00:00, 983.06it/s]
```

LogReg finished.
ROC-AUC: 0.9512987012987013

```
# =====
# MLflow Experiment Tracking : RandomForest
# =====

from pathlib import Path
import numpy as np
import mlflow
import mlflow.sklearn

from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    roc_auc_score, accuracy_score, precision_score, recall_score,
    f1_score, confusion_matrix
)

# -----
# 0) MLflow local tracking (portable, no hardcoded path)
# -----
PROJECT_ROOT = Path.cwd().resolve()
MLRUNS_DIR = PROJECT_ROOT / "mlruns"
(MLRUNS_DIR / ".trash").mkdir(parents=True, exist_ok=True)

mlflow.set_tracking_uri(MLRUNS_DIR.as_uri())
EXPERIMENT_NAME = "ravindra_cardio_risk_mllops_assignment1"
```

```

mlflow.set_experiment(EXPERIMENT_NAME)

# -----
# 1) Stop any active run (re-run safe)
# -----
if mlflow.active_run() is not None:
    mlflow.end_run(status="KILLED")

# -----
# 2) Build preprocessing + model pipeline
# (RF doesn't need scaling)
# -----
num_pipe = Pipeline(steps=[
    ("impute", SimpleImputer(strategy="median"))
])

cat_pipe = Pipeline(steps=[
    ("impute", SimpleImputer(strategy="most_frequent")),
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])
preprocessor = ColumnTransformer(
    transformers=[
        ("num", num_pipe, num_cols),
        ("cat", cat_pipe, cat_cols),
    ],
    remainder="drop"
)

rf_params = dict(
    n_estimators=400,
    max_depth=None,
    min_samples_split=2,
    min_samples_leaf=1,
    max_features="sqrt",
    random_state=42,
    n_jobs=-1
)

rf_pipeline = Pipeline(steps=[
    ("prep", preprocessor),
    ("clf", RandomForestClassifier(**rf_params))
])

# -----
# 3) Train + log to MLflow
# -----
run = mlflow.start_run(run_name="tuned_random_forest_notebook")

try:

```

```

mlflow.set_tag("model_family", "random_forest")
mlflow.set_tag("dataset", "UCI Heart Disease (Cleveland)")

for k, v in rf_params.items():
    mlflow.log_param(k, v)

rf_pipeline.fit(X_train, y_train)

y_prob = rf_pipeline.predict_proba(X_test)[:, 1]
y_pred = (y_prob >= 0.5).astype(int)

auc = roc_auc_score(y_test, y_prob)
acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred, zero_division=0)
rec = recall_score(y_test, y_pred, zero_division=0)
f1 = f1_score(y_test, y_pred, zero_division=0)
cm = confusion_matrix(y_test, y_pred)

mlflow.log_metric("roc_auc", float(auc))
mlflow.log_metric("accuracy", float(acc))
mlflow.log_metric("precision", float(prec))
mlflow.log_metric("recall", float(rec))
mlflow.log_metric("f1", float(f1))

mlflow.log_param("cm_00", int(cm[0, 0]))
mlflow.log_param("cm_01", int(cm[0, 1]))
mlflow.log_param("cm_10", int(cm[1, 0]))
mlflow.log_param("cm_11", int(cm[1, 1]))

mlflow.sklearn.log_model(
    sk_model=rf_pipeline,
    name="model",
    input_example=X_test.iloc[:3]
)

print("RandomForest finished.")
print("ROC-AUC:", auc)

except Exception:
    mlflow.end_run(status="FAILED")
    raise

finally:
    if mlflow.active_run() is not None:
        mlflow.end_run(status="FINISHED")

C:\Users\Ravindra\AppData\Roaming\Python\Python313\site-packages\
mlflow\utils.py:452: UserWarning: Hint: Inferred schema contains
integer column(s). Integer columns in Python cannot represent missing
values. If your input data contains missing values at inference time,

```

```
it will be encoded as floats and will cause a schema enforcement
error. The best way to avoid this problem is to infer the model schema
based on a realistic data sample (training dataset) that includes
missing values. Alternatively, you can declare integer columns as
doubles (float64) whenever these columns may have missing values. See
`Handling Integers With Missing Values
<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-missing-values>`_ for more details.
    warnings.warn(
Downloading artifacts: 100%|██████████| 7/7
[00:00<00:00, 742.19it/s]

RandomForest finished.
ROC-AUC: 0.9583333333333333

#mlflow run

#!/python -m mlflow ui --backend-store-uri "../mlruns" --port 5000
# stop run

import mlflow

# Close any possibly-open run
try:
    mlflow.end_run(status="KILLED")
except Exception:
    pass

# Ensure no active run remains
while mlflow.active_run() is not None:
    mlflow.end_run(status="KILLED")

print("active_run:", mlflow.active_run()) # must print None
active_run: None

# Experiment tracking RandomForest

import itertools
from pathlib import Path
import pandas as pd

import mlflow
import mlflow.sklearn
from mlflow.models.signature import infer_signature

from sklearn.pipeline import Pipeline
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score,
```

```

recall_score, roc_auc_score

# --- Notebook-safe cleanup (prevents "run already active") ---
try:
    mlflow.end_run(status="KILLED")
except Exception:
    pass

# --- MLflow setup ---
PROJECT_ROOT = Path(".").resolve()
MLRUNS_DIR = (PROJECT_ROOT / "mlruns").resolve()
mlflow.set_tracking_uri(MLRUNS_DIR.as_uri())
mlflow.set_experiment("ravindra_cardio_risk_single_run_tuning")

ART_DIR = (PROJECT_ROOT / "artifacts" / "tuning").resolve()
ART_DIR.mkdir(parents=True, exist_ok=True)

# --- Hyperparameter grid ---
grid = {
    "n_estimators": [150, 300],
    "max_depth": [None, 6, 10],
    "min_samples_leaf": [1, 2],
}

def compute_metrics(y_true, y_pred, y_prob):
    return {
        "accuracy": float(accuracy_score(y_true, y_pred)),
        "precision": float(precision_score(y_true, y_pred,
zero_division=0)),
        "recall": float(recall_score(y_true, y_pred,
zero_division=0)),
        "roc_auc": float(roc_auc_score(y_true, y_prob)),
    }

# --- Single MLflow run for the whole tuning session ---
with mlflow.start_run(run_name="rf_hparam_tuning_single_run"):
    mlflow.set_tag("course", "ML Ops Assignment 1")
    mlflow.set_tag("dataset", "UCI Heart Disease (Cleveland)")
    mlflow.log_param("model_family", "RandomForestClassifier")
    mlflow.log_param("threshold_used", 0.5)
    mlflow.log_param("grid_size",
len(list(itertools.product(*grid.values()))))

    keys = list(grid.keys())
    combos = list(itertools.product(*[grid[k] for k in keys]))

    results = []
    best_model = None
    best_row = None
    best_score = -1.0

```

```

for trial_id, values in enumerate(combos, start=1):
    params = dict(zip(keys, values))

    model = Pipeline(steps=[
        ("prep", preprocessor),
        ("clf", RandomForestClassifier(random_state=42, **params))
    ])

    model.fit(X_train, y_train)
    prob = model.predict_proba(X_test)[:, 1]
    pred = (prob >= 0.5).astype(int)

    m = compute_metrics(y_test, pred, prob)

    row = {"trial_id": trial_id, **params, **m}
    results.append(row)

    # Track progress in MLflow as "trial metrics" (same run)
    mlflow.log_metric("trial_roc_auc", m["roc_auc"],
step=trial_id)
    mlflow.log_metric("trial_accuracy", m["accuracy"],
step=trial_id)

    # Select best by ROC-AUC (change if needed)
    if m["roc_auc"] > best_score:
        best_score = m["roc_auc"]
        best_model = model
        best_row = row

    # --- Save & log the full tuning table ---
    results_df = pd.DataFrame(results).sort_values("roc_auc",
ascending=False)
    csv_path = ART_DIR / "rf_hparam_trials.csv"
    results_df.to_csv(csv_path, index=False)
    mlflow.log_artifact(str(csv_path), artifact_path="tuning")

    # --- Log the best params + best metrics as top-level outputs ---
    best_params = {k: best_row[k] for k in keys}
    mlflow.log_params({f"best_{k)": v for k, v in
best_params.items()})

    mlflow.log_metric("best_roc_auc", float(best_row["roc_auc"]))
    mlflow.log_metric("best_accuracy", float(best_row["accuracy"]))
    mlflow.log_metric("best_precision", float(best_row["precision"]))
    mlflow.log_metric("best_recall", float(best_row["recall"]))

    # --- Log the best model pipeline ---
    signature = infer_signature(X_test,
best_model.predict_proba(X_test))

```

```

mlflow.sklearn.log_model(
    sk_model=best_model,
    artifact_path="best_model",
    signature=signature,
    input_example=X_test.iloc[:3]
)
BEST_RUN_ID = mlflow.active_run().info.run_id
MODEL_NAME = "ravindra_cardio_risk_model"

model_uri = f"runs:{BEST_RUN_ID}/best_model"

mlflow.register_model(
    model_uri=model_uri,
    name=MODEL_NAME
)
print("Best model registered successfully")

print("Done. Trials logged in one run. Best ROC-AUC:", best_score)

2025/12/26 23:43:04 WARNING mlflow.models.model: `artifact_path` is
deprecated. Please use `name` instead.
Downloading artifacts: 100%|██████████| 7/7
[00:00<00:00, 643.11it/s]
C:\Users\Ravindra\AppData\Roaming\Python\Python313\site-packages\
mlflow\tracking\_model_registry\utils.py:215: FutureWarning: The
filesystem model registry backend (e.g., './mlruns') will be
deprecated in February 2026. Consider transitioning to a database
backend (e.g., 'sqlite:///mlflow.db') to take advantage of the latest
MLflow features. See https://github.com/mlflow/mlflow/issues/18534 for
more details and migration guidance. For migrating existing data,
https://github.com/mlflow/mlflow-export-import can be used.
    return FileStore(store_uri)
Successfully registered model 'ravindra_cardio_risk_model'.
2025/12/26 23:43:08 WARNING mlflow.tracking._model_registry.fluent:
Run with id 2cfceb8e3e5c4f8d9d4357c12d203ad9 has no artifacts at
artifact path 'best_model', registering model based on models:/m-
1ab57f3548e84d0593b5d8019af157c4 instead

Best model registered successfully
Done. Trials logged in one run. Best ROC-AUC: 0.9664502164502164
Created version '1' of model 'ravindra_cardio_risk_model'.

```

TASK 4 : Model Packaging & Reproducibility

```

# Export final model

import numpy as np
from sklearn.metrics import roc_auc_score

```

```

def _is_predictor(obj):
    return hasattr(obj, "predict") and (hasattr(obj, "predict_proba")
or hasattr(obj, "decision_function"))

def _predict_scores(obj, X):
    # Prefer predict_proba for binary classification
    if hasattr(obj, "predict_proba"):
        p = obj.predict_proba(X)
        # take positive class if 2D
        if isinstance(p, np.ndarray) and p.ndim == 2 and p.shape[1] >=
2:
            return p[:, 1]
        return p
    # fallback: decision_function
    return obj.decision_function(X)

def select_final_pipeline_from_globals(X_eval, y_eval,
metric="roc_auc"):
    best_name, best_obj, best_score = None, None, -1.0
    tried = 0

    for name, obj in globals().items():
        # skip obvious non-model objects
        if name.startswith("_"):
            continue
        if not _is_predictor(obj):
            continue

        # Try to compute AUC; skip objects that can't score
        try:
            scores = _predict_scores(obj, X_eval)
            # Ensure 1D
            scores = np.asarray(scores).reshape(-1)
            score = roc_auc_score(y_eval, scores)
            tried += 1
        except Exception:
            continue

        if score > best_score:
            best_name, best_obj, best_score = name, obj, score

    if best_obj is None:
        raise RuntimeError(
            "No trained model/pipeline found in globals(). "
            "Train a model first (Pipeline.fit) and ensure it exposes
predict_proba or decision_function."
    )

    return best_name, best_obj, float(best_score), tried

```

```
best_name, final_pipeline, best_auc, tried =
select_final_pipeline_from_globals(X_test, y_test)

summary = (
    "Auto-selection complete.\n"
    "All trained models available in the notebook were evaluated using
ROC-AUC on the test set.\n"
    f"Selected model: {best_name}\n"
    f"Test ROC-AUC: {best_auc:.4f}\n"
    f"Candidate models evaluated: {tried}\n"
    "Pipeline structure: prep (preprocessing) → clf (classifier)\n"
    "Because preprocessing and prediction are bundled into a single
pipeline, "
    "this model is ready for packaging and reproducible inference."
)

print(summary)

Auto-selection complete.
All trained models available in the notebook were evaluated using ROC-
AUC on the test set.
Selected model: logreg
Test ROC-AUC: 0.9665
Candidate models evaluated: 8
Pipeline structure: prep (preprocessing) → clf (classifier)
Because preprocessing and prediction are bundled into a single
pipeline, this model is ready for packaging and reproducible
inference.

# -----
# TASK 4 FINAL: Clean requirements.txt
# -----
from pathlib import Path
import sys
import re
import pkgutil
import importlib

# 1) Resolve project root portably (no hard-coded paths)
PROJECT_ROOT = Path.cwd().resolve()

# If notebook is inside a subfolder, try walking up until we find a
# marker folder
# (choose either "mlruns" or "artifacts" if you have them)
for _ in range(5):
    if (PROJECT_ROOT / "mlruns").exists() or (PROJECT_ROOT /
"artifacts").exists():
        break
```

```

PROJECT_ROOT = PROJECT_ROOT.parent

REQ_PATH = PROJECT_ROOT / "requirements.txt"

# 2) Helper: safe version lookup (works on Py 3.13)
try:
    from importlib.metadata import version as pkg_version
except Exception:
    pkg_version = None

def _ver(pkg_name: str) -> str | None:
    if pkg_version is None:
        return None
    try:
        return pkg_version(pkg_name)
    except Exception:
        return None

# 3) Define a CLEAN, minimal set (edit only if you truly used more)
# These are the typical packages for your notebook: EDA + sklearn
models + MLflow + API
base_packages = [
    "numpy",
    "pandas",
    "scikit-learn",
    "matplotlib",
    "joblib",
    "mlflow",
]
# Optional (only include if you actually used them in code)
optional_packages = [
    "seaborn",
    "scipy",
    "fastapi",
    "uvicorn",
    "pydantic",
]
# 4) Keep only installed packages (prevents pip failures on graders'
machine)
clean = []
for p in base_packages + optional_packages:
    v = _ver(p)
    if v:
        clean.append((p, v))

# 5) Write requirements.txt with pinned versions (clean + readable)
lines = [
    "# Auto-generated for reproducibility (Task 4 - Model Packaging &

```

```

Reproducibility)",
    f"# Python: {sys.version.split()[0]}",
    "",
]
for p, v in clean:
    lines.append(f"{p}=={v}")

REQ_PATH.write_text("\n".join(lines) + "\n", encoding="utf-8")

print("requirements.txt created at:", REQ_PATH)
print("\n--- requirements.txt (preview) ---")
print(REQ_PATH.read_text(encoding="utf-8"))

# 6) (Optional but recommended) Log it to MLflow if a run is active
# If no run is active, we still created the file which satisfies the
rubric.
try:
    import mlflow
    if mlflow.active_run() is not None:
        mlflow.log_artifact(str(REQ_PATH),
artifact_path="reproducibility")
        print("\nLogged requirements.txt to MLflow under:
reproducibility/")
    else:
        print("\nNo active MLflow run. File is created locally ")
except Exception as e:
    print("\nMLflow logging skipped due to:", repr(e))

requirements.txt created at: C:\Users\Ravindra\Documents\
cardio_risk_ops\notebooks\requirements.txt

--- requirements.txt (preview) ---
# Auto-generated for reproducibility (Task 4 - Model Packaging &
Reproducibility)
# Python: 3.13.7

numpy==2.3.5
pandas==2.3.3
scikit-learn==1.8.0
matplotlib==3.10.7
joblib==1.5.2
mlflow==3.8.0
seaborn==0.13.2
scipy==1.16.3
fastapi==0.127.0
uvicorn==0.40.0
pydantic==2.12.5

No active MLflow run. File is created locally

```

